

THE-CERT-SCHOOL

FREE PRACTICE TEST/QUIZ/EXAM FOR
AWS CSA/AWS CDA/AWS SYSOPS

[HOME](#)[CONTACT](#)

Protected: Quiz7

By: cert-school-admin / On: December 15, 2016 / In:


AWS, AWS CSA / With: 6 Comments /

Results

0 of 30 questions answered correctly

Your time: 00:00:02

You have reached 0 of 30 points, (0%)



Categories

Data Security	0%
Designing highly available, cost-efficient, fault-tolerant, scalable	0%
Implementation/Deployment	0%
Troubleshooting	0%

**ALL COURSES 10\$ @
WWW.UDEMY.COM**

Promocode: **BREAKTHRU**



MEGA 10\$ OFFER @UDEMY

Start Date: 2/22/2017

11:59PM PST

End Date: 3/1/2017 6:00AM

PST

Promocode: **BREAKTHRU**

Discount: \$10

**Want to pass the AWS
Solutions Architect -
Associate Exam? Do this
course!**



Thank you for taking this mock exam. Unfortunately, you were unsuccessful in this attempt.

[View questions](#)

1	2	3	4	5	6	7	8	9	10	11
12	13	14	15	16	17	18	19	20	21	22

Answered Review

1. Question

Your customer is willing to consolidate their log streams (access logs, application logs, security logs etc.) in one single system. Once consolidated, the customer wants to analyse these logs in real time based on heuristics. From time to time, the customer needs to validate heuristics, which requires going back to data samples extracted from the last 12 hours? What is the best approach to meet your customer's requirements?

- ☐ Configure Amazon Cloud Trail to receive custom logs, use EMR to apply heuristics the logs
- ☐ Setup an Auto Scaling group of EC2 syslogd servers, store the logs on S3 use EMR to apply heuristics on the logs
- ☐ Send all the log events to Amazon SQS. Setup an Auto Scaling group of EC2 servers to consume the logs and apply the heuristics.
- ☐ **Send all the log events to Amazon Kinesis develop a client process to apply heuristics on the logs**

Incorrect

Want to pass the AWS Certified Developer - Associate 2017? Do this course!



Want to pass the AWS Certified SysOps Administrator - Associate 2017? Do this course!



Master AWS CodeDeploy



Learn AWS The Hard Way



DevOps: CI/CD using AWS CodePipeline & Elastic Beanstalk

Use Amazon Kinesis Streams to collect and process large streams of data records in real time. You'll create data-processing applications, known as Amazon Kinesis Streams applications. A typical Amazon Kinesis Streams application reads data from an Amazon Kinesis stream as data records. These applications can use the Amazon Kinesis Client Library, and they can run on Amazon EC2 instances. The processed records can be sent to dashboards, used to generate alerts, dynamically change pricing and advertising strategies, or send data to a variety of other AWS services.



TESTIMONIALS

**Alexander
Korolyov**

Thank you the cert school! Your exam quizzes really helpful during the preparation.
Thank you again and continue to do the same!

2. Question

What is an isolated database environment running in the cloud (Amazon RDS) called?

- ☒ DB Instance
- ☐ DB Server
- ☐ DB Unit
- ☐ DB Volume

Incorrect

A DB instance is an isolated database environment running in the cloud. It is the basic building block of Amazon RDS. A DB instance can contain multiple user-created databases, and can be accessed using the same client tools and applications you might use to access a stand-alone database instance. DB instances are simple to create and modify with the Amazon AWS command line tools, Amazon RDS APIs, or the AWS Management RDS Console.

TAGS

AWS AWS CDA AWS
CSA AWS SysOps

Lessons

3. Question

IAM provides several policy templates you can use to automatically assign permissions to the groups you create. The ____ policy template gives the Admin group permission to access all account resources, except your AWS account information

- ☐ **Power User Access**
- ☐ Read Only Access
- ☐ Read Only Access
- ☐ Administrator Access

Incorrect

The power user role provides an AWS Directory Service user or group with full access to AWS services and resources, but does not allow management of IAM users and groups. The following is the policy for this role.

```
{  
  "Version" : "2012-10-17",  
  "Statement" : [  
    {  
      "Effect" : "Allow",  
      "NotAction" : "iam:*",  
      "Resource" : "*" }  
  ]  
}
```

4. Question

While performing the volume status checks, if the status is insufficient-data, what does it mean?

- ☐ the checks is not yet started
- ☐ the check has passed

- ☐ the check has failed
- ☐ **the checks may still be in progress on the volume**

Incorrect

Volume status checks are automated tests that run every 5 minutes and return a pass or fail status. If all checks pass, the status of the volume is ok. If a check fails, the status of the volume is impaired. If the status is insufficient-data, the checks may still be in progress on the volume. You can view the results of volume status checks to identify any impaired volumes and take any necessary actions.

5. Question

Which of the below policy provides only full access to Amazon S3 services and resources?

- ☐ { "Version": "2016-10-17", "Statement": [{ "Effect": "Allow", "Action": "*", "Resource": "*" }] }
- ☐ { "Version": "2016-10-17", "Statement": [{ "Effect": "Allow", "Action": "*:s3:*", "Resource": "*" }] }
- ☐ **{ "Version": "2016-10-17", "Statement": [{ "Effect": "Allow", "Action": "s3:*", "Resource": "*" }] }**
- ☐ { "Version": "2016-10-17", "Statement": [{ "Effect": "Allow", "Action": "*:s3", "Resource": "*" }] }

Incorrect

Refer:

http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_elements.html

6. Question

You have an EC2 Security Group with several running EC2 instances. You change the Security Group rules to allow inbound traffic on a new port and protocol, and launch several new instances in the same Security Group. The new rules apply:

- ☒ **Immediately to all instances in the security group.**
- ☐ Immediately to the new instances, but old instances must be stopped and restarted before the new rules apply.
- ☐ To all instances, but it may take several minutes for old instances to see the changes.
- ☐ Immediately to the new instances only.

Incorrect

When you add a rule to a security group, the new rule is automatically applied to any instances associated with the security group. You can assign a security group to an instance when you launch the instance. When you add or remove rules, those changes are automatically applied to all instances to which you've assigned the security group.

7. Question

When automatic failover occurs, Amazon RDS will emit a DB Instance event to inform you that automatic failover occurred. You can use the ____ to return information about events related to your DB Instance

- ☒ **DescribeEvents**
- ☐ FetchFailure
- ☐ FetchEvents
- ☐ DescribeFailure

Incorrect

Amazon RDS will emit a DB Instance event to inform you that automatic failover occurred. You can use the DescribeEvents to return information about events related to your DB Instance, or click the “DB Events” section of the AWS Management Console

8. Question

Which of the below is true about S3 Cross Region Replication(Select 2 answers)?

- ☐ Cross-region replication requires that versioning is enabled only in source bucket and not needed in destination bucket
- ☐ By activating cross-region replication, actions performed by lifecycle configuration will be also replicated
- ☐ **Cross-region replication is the automatic, asynchronous copying of objects across buckets in different AWS regions**
- ☐ **By activating cross-region replication, Amazon S3 will replicate newly created objects, object updates, and object deletions from a source bucket into a destination bucket in a different region**
- ☐ Amazon S3 console allows you to delete cross-region replication

Incorrect

Cross-region replication is the automatic, asynchronous copying of objects across buckets in different AWS regions. By activating cross-region replication, Amazon S3 will replicate newly created objects, object updates, and object deletions from a source bucket into a destination bucket in a different region. Cross-region replication has specific

requirements that define what can and cannot be replicated across regions based on how the object is created and how it is encrypted.

Refer: <http://docs.aws.amazon.com/AmazonS3/latest/UG/cross-region-replication.html>

9. Question

Select the correct set of steps for exposing the snapshot only to specific AWS accounts

- ☐ SelectPublic, enter the IDs of those AWS accounts, and clickSave.
- ☐ Select public for all the accounts and check mark those accounts with whom you want to expose the snapshots and click save.
- ☐ **SelectPrivate, enter the IDs of those AWS accounts, and clickSave**
- ☐ SelectPublic, mark the IDs of those AWS accounts as private, and clickSave.

Incorrect

To expose the snapshot to only specific AWS accounts, choose Private, enter the ID of the AWS account (without hyphens) in the AWS Account Number field, and choose Add Permission. Repeat until you've added all the required AWS accounts.

10. Question

What does the "Server Side Encryption" option on Amazon S3 provide?

- ☐ It provides an encrypted virtual disk in the Cloud.
- ☐ **It encrypts the files that you send to Amazon S3, on the server side.**
- ☐ It doesn't exist for Amazon S3, but only for Amazon EC2.
- ☐ It allows to upload files using an SSL endpoint, for a secure transfer.

Incorrect

Server-side encryption is about protecting data at rest. Server-side encryption with Amazon S3-managed encryption keys (SSE-S3) employs strong multi-factor encryption. Amazon S3 encrypts each object with a unique key. As an additional safeguard, it encrypts the key itself with a master key that it regularly rotates. Amazon S3 server-side encryption uses one of the strongest block ciphers available, 256-bit Advanced Encryption Standard (AES-256), to encrypt your data. Amazon S3 supports bucket policies that you can use if you require server-side encryption for all objects that are stored in your bucket.

11. Question

A company is building a two-tier web application to serve dynamic transaction-based content. The data tier is leveraging an Online Transactional Processing(OLTP)database. What services should you leverage to enable an elastic and scalable web tier?

- ☐ AmazonEC2,Amazon Dynamo DB, and Amazon S3
- ☐ Elastic Load Balancing, Amazon EC2, and Amazon RDS
- ☐ Amazon RDS with Multi-AZ and Auto Scaling
- ☐ **Elastic Load Balancing , Amazon EC2, and Auto Scaling**

Incorrect

Elastic and Scalable are the keywords.

Auto Scaling helps you ensure that you have the correct number of Amazon EC2 instances available to handle the load for your application.

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic.

12. Question

What does Amazon Cloud Formation provide?

- ☐ The ability to setup Auto scaling for Amazon EC2 instances
- ☒ **Create templates for the service or application architectures**
- ☐ A template to map network resources for Amazon Web Services
- ☐ A container for Amazon Services

Incorrect

AWS CloudFormation simplifies provisioning and management on AWS. You can create templates for the service or application architectures you want and have AWS CloudFormation use those templates for quick and reliable provisioning of the services or applications (called "stacks"). You can also easily update or replicate the stacks as needed. This collection of sample templates will help you

get started with AWS Cloud Formation and quickly build your own templates.

13. Question

How are the EBS snapshots saved on Amazon S3?

- ☐ Exponentially
- ☐ EBS snapshots are not stored in the Amazon S3
- ☐ **Incrementally**
- ☐ Are expressly prohibited under all circumstances.
- ☐ Decrementally

Incorrect

You can back up the data on your EBS volumes to Amazon S3 by taking point-in-time snapshots. Snapshots are incremental backups, which means that only the blocks on the device that have changed after your most recent snapshot are saved. This minimizes the time required to create the snapshot and saves on storage costs. When you delete a snapshot, only the data unique to that snapshot is removed. Active snapshots contain all of the information needed to restore your data (from the time the snapshot was taken) to a new EBS volume.

14. Question

When creation of an EBS snapshot is initiated, but not completed, the EBS volume:

- ☐ **Can be used while the snapshot is in progress**
- ☐ Cannot be detached or attached to an EC2 instance until the snapshot completes

- ☐ Can be used but there should be a delay in IO operations
- ☐ Can be used in read-only mode while the snapshot is in progress

Incorrect

Snapshots occur asynchronously; the point-in-time snapshot is created immediately, but the status of the snapshot is pending until the snapshot is complete (when all of the modified blocks have been transferred to Amazon S3), which can take several hours for large initial snapshots or subsequent snapshots where many blocks have changed. While it is completing, an in-progress snapshot is not affected by ongoing reads and writes to the volume.

15. Question

To help you manage your Amazon EC2 instances, images, and other Amazon EC2 resources, you can assign your own meta data to each resource in the form of _____

- ☐ functions
- ☒ tags
- ☐ special filters
- ☐ wildcards

Incorrect**Tagging Your Amazon EC2 Resources**

To help you manage your instances, images, and other Amazon EC2 resources, you can optionally assign your own metadata to each resource in the form of tags. Tags enable you to categorize your AWS resources in different ways, for

example, by purpose, owner, or environment. This is useful when you have many resources of the same type.

16. Question

Regarding the attaching of ENI to an instance, what does 'warm attach' refer to?

- ☐ Attaching an ENI to an instance when it is starting
- ☐ Attaching an ENI to an instance when it is running
- ☐ **Attaching an ENI to an instance when it is stopped.**
- ☐ Attaching an ENI to an instance during the launch process

Incorrect

You can attach an elastic network interface to an instance when it's running (hot attach), when it's stopped (warm attach), or when the instance is being launched (cold attach).

17. Question

You need to configure an Amazon S3 bucket to serve static assets for your public-facing web application. Which methods ensure that all objects uploaded to the bucket are set to public read? Choose 2 answers

- ☐ **Set permissions on the object to public read during upload**
- ☐ Amazon S3 objects default to public read, so no action is needed
- ☐ Configure the bucket ACL to set all objects to public read
- ☐ **Configure the bucket policy to set all objects to public read**

- ☐ Use AWS Identity and Access Management roles to set the bucket to public read

Incorrect

You can use ACLs to grant permissions to individual AWS accounts; however, it is strongly recommended that you do not grant public access to your bucket using an ACL. So the recommended approach is create bucket policy, but not ACL. You must grant read permission on the specific objects to make them publicly accessible so that your users can view them on your website. You make objects publicly readable by using either the object ACL or by writing a bucket policy

18. Question

You can use ____ and ____ to help secure the instances in your VPC.

- ☐ **security groups and network ACLs**
- ☐ security groups and multi-factor authentication
- ☐ security groups and biometric authentication
- ☐ NCLs and 2-Factor authentication

Incorrect

Security groups — Act as a firewall for associated Amazon EC2 instances, controlling both inbound and outbound traffic at the instance level

Network access control lists (ACLs) — Act as a firewall for associated subnets, controlling both inbound and outbound traffic at the subnet level

19. Question

Every user you create in the IAM system starts with _____

- ☐ Partial permissions
- ☐ Power user permissions
- ☐ Full permissions
- ☐ **No permissions**

Incorrect

Permissions let you specify who has access to AWS resources, and what actions they can perform on those resources. Every IAM user starts with no permissions. In other words, by default, users can do nothing, not even view their own access keys. To give a user permission to do something, you can add the permission to the user (that is, attach a policy to the user) or add the user to a group that has the desired permission.

20. Question

In RDS instance, you must increase storage size in increments of at least _____ %

- ☐ 15
- ☐ **10**
- ☐ 20
- ☐ 30

Incorrect

AllocatedStorage: The new storage capacity of the RDS instance. Changing this setting does not result in an outage

and the change is applied during the next maintenance window unless ApplyImmediately is set to true for this request.

Constraints: Value supplied must be at least 10% greater than the current value. Values that are not at least 10% greater than the existing value are rounded up so that they are 10% greater than the current value.

21. Question

What does the AWS Storage Gateway provide?

- ☐ It provides an encrypted SSL endpoint for backups in the Cloud.
- ☐ A direct encrypted connection to Amazon S3
- ☐ It's a backup solution that provides an on-premises Cloud storage
- ☐ **It allows to integrate on-premises IT environments with Cloud Storage.**

Incorrect

AWS Storage Gateway connects an on-premises software appliance with cloud-based storage to provide seamless integration with data security features between your on-premises IT environment and the Amazon Web Services (AWS) storage infrastructure. You can use the service to store data in the AWS cloud for scalable and cost-effective storage that helps maintain data security

22. Question

Resources that are created in AWS are identified by a unique identifier called an _____

- ☐ Amazon Resource Name tag
- ☐ **Amazon Resource Name**
- ☐ Amazon Resource Namespace
- ☐ Amazon Resource Number

Incorrect

Amazon Resource Names (ARNs) uniquely identify AWS resources. We require an ARN when you need to specify a resource unambiguously across all of AWS, such as in IAM policies, Amazon Relational Database Service (Amazon RDS) tags, and API calls.

23. Question

What does Amazon EC2 provide?

- ☐ **Virtual servers in the Cloud.**
- ☐ Physical servers, remotely managed by the customer.
- ☐ A platform to run code (Java, PHP, Python), paying on an hourly basis.
- ☐ Computer Clusters in the Cloud

Incorrect

Amazon Elastic Compute Cloud (Amazon EC2) provides scalable computing capacity in the Amazon Web Services (AWS) cloud. Using Amazon EC2 eliminates your need to invest in hardware up front, so you can develop and deploy applications faster. You can use Amazon EC2 to launch as many or as few virtual servers as you need, configure security and networking, and manage storage. Amazon EC2 enables you to scale up or down to handle changes in

requirements or spikes in popularity, reducing your need to forecast traffic.

24. Question

Which services allow the customer to retain full administrative privileges of the underlying EC2 instances? Choose 2 answers

- ☐ **Amazon Elastic Map Reduce**
- ☐ Amazon Relational Database Service
- ☐ Amazon ElastiCache
- ☐ **AWS Elastic Beanstalk**
- ☐ Amazon DynamoDB

Incorrect

- **AWS provides the root or system privileges only for a limited set of services, which includes**
 - **Elastic Cloud Compute (EC2)**
 - **Elastic MapReduce (EMR)**
 - **Elastic BeanStalk**
 - **Opswork**
- **AWS does not provide root privileges for managed services like RDS, DynamoDB, S3, Glacier etc**

25. Question

What type of block cipher does Amazon S3 offer for server side encryption?

- ☐ Triple DES
- ☐ Blowfish
- ☐ **Advanced Encryption Standard**

☐ RC5

Incorrect

Server-side encryption is about protecting data at rest. Server-side encryption with Amazon S3-managed encryption keys (SSE-S3) employs strong multi-factor encryption. Amazon S3 encrypts each object with a unique key. As an additional safeguard, it encrypts the key itself with a master key that it regularly rotates. Amazon S3 server-side encryption uses one of the strongest block ciphers available, 256-bit Advanced Encryption Standard (AES-256), to encrypt your data.

26. Question

You manually launch a NAT AMI in a public subnet. The network is properly configured. Security groups and network access control lists are property configured. Instances in a private subnet can access the NAT. The NAT can access the Internet. However, private instances cannot access the Internet. What additional step is required to allow access from the private instances?

- ☐ Disable Source/Destination Check on the private instances.
- ☐ Enable Source/Destination Check on the NAT instance
- ☐ Enable Source/Destination Check on the private Instances.
- ☐ **Disable Source/Destination Check on the NAT instance.**

Incorrect

Each EC2 instance performs source/destination checks by default. This means that the instance must be the source or destination of any traffic it sends or receives. However, a NAT instance must be able to send and receive traffic when the source or destination is not itself. Therefore, you must

disable source/destination checks on the NAT instance. You can disable the SrcDestCheck attribute for a NAT instance that's either running or stopped using the console or the command line.

To disable source/destination checking using the console

1. Open the Amazon EC2 console.

2. In the navigation pane, choose Instances.

3. Select the NAT instance, choose Actions, select Networking, and then select Change Source/Dest. Check.

4. For the NAT instance, verify that this attribute is disabled. Otherwise, choose Yes, Disable.

27. Question

A t2.medium EC2 instance type must be launched with what type of Amazon Machine Image (AMI)?

- ☐ An Instance store Paravirtual AMI
- ☒ **An Amazon EBS-backed Hardware Virtual Machine AMI**
- ☐ An Instance store Hardware Virtual Machine AMI
- ☐ An Amazon EBS-backed Paravirtual AMI

Incorrect

Refer: <https://aws.amazon.com/amazon-linux-ami/instance-type-matrix/>

28. Question

Select the correct statement for Amazon RedShift Vs Amazon EMR (choose 2 answers)

- ☐ **Amazon Redshift is ideal for large volumes of structured data that you want to persist**

☐ **Amazon EMR is ideal for processing and transforming unstructured or semi-structured data to bring in to Amazon Redshift**

☐ Amazon Redshift is ideal for processing and transforming unstructured or semi-structured data to bring in to Amazon EMR

☐ Amazon EMR is ideal for large volumes of structured data that you want to persist

☐ Amazon Redshift is much better option for data sets that are relatively transitory, not stored for long-term use

Incorrect

Amazon Redshift is ideal for large volumes of structured data that you want to persist and query using standard SQL and your existing BI tools. Amazon EMR is ideal for processing and transforming unstructured or semi-structured data to bring in to Amazon Redshift and is also a much better option for data sets that are relatively transitory, not stored for long-term use.

29. Question

A customer has a single 3-TB volume on-premises that is used to hold a large repository of images and print layout files. This repository is growing at 500 GB a year and must be presented as a single logical volume. The customer is becoming increasingly constrained with their local storage capacity and wants an off-site backup of this data, while maintaining low-latency access to their frequently accessed data. Which AWS Storage Gateway configuration meets the customer requirements?

☐ Gateway-Stored volumes with snapshots scheduled to Amazon S3

- ☐ Gateway-Virtual Tape Library with snapshots to Amazon Glacier
- ☐ **Gateway-Cached volumes with snapshots scheduled to Amazon S3**
- ☐ Gateway-Virtual Tape Library with snapshots to Amazon S3

Incorrect

Gateway-cached volumes allow you to utilize Amazon S3 for your primary data, while retaining some portion of it locally in a cache for frequently accessed data. These volumes minimize the need to scale your on-premises storage infrastructure, while still providing your applications with low-latency access to their frequently accessed data. You can create up to 32 volumes up to 32 TB in size each, for a total of 1 PB of data capacity per gateway, and mount them as iSCSI devices from your on-premises application servers. Data written to these volumes is stored in Amazon S3, with only a cache of recently written and recently read data stored locally on your on-premises storage hardware.

30. Question

In the Amazon cloud watch, which metric should I be checking to ensure that your DB Instance has enough free storage space?

- ☐ **FreeStorageSpace**
- ☐ FreeStorage
- ☐ FreeDBStorageSpace
- ☐ FreeStorageVolume

Incorrect

Amazon Relational Database Service sends metrics to CloudWatch for each active database instance every minute. Detailed monitoring is enabled by default.
FreeStorageSpace:The amount of available storage space.

Units: Bytes

Related

Quiz2

January 5, 2017

In "AWS"

Passed the AWS
CSA-A exam with
85%

Author - Yuvaraj
Pyneni First thing,
exam now has 60
questions & 80
Minutes, I have
December 19, 2016
In "Lessons"

Passed my AWS-
CSA-Associate
Exam with 90%
score

I have passed my
CSA-A exam last
week with 90%
score after tons of
study and sitting
December 5, 2016
In "Lessons"

Previous Post: [Passed my AWS-CSA-Associate Exam with 90% score](#)

Next Post: [Passed the AWS CSA-A exam with 85%](#)

6 COMMENTS

Questions are good, but the website has too many ads which are crashing my browser. Whatever link I click it simply open lot of windows and it sucks

[HADOOPER](#) | [3 MONTHS AGO](#) | [PERMALINK](#) | [REPLY](#)

AUTHOR Hello Sir,

I have corrected it. Thanks for your feedback

Admin

CERT-SCHOOL-ADMIN |

3 MONTHS AGO | [PERMALINK](#) | [REPLY](#)

<http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/rds-metricscollected.html>

Q30 – the answer is wrong.

FreeStorageSpace is for storage space, not random access memory.

FreeableMemory

The amount of available random access memory.

Units: Bytes

FreeStorageSpace

The amount of available storage space.

Units: Bytes

SIMON | 1 MONTH AGO | [PERMALINK](#) | [REPLY](#)

AUTHOR Yes sir,

Updated it. Thanks for your feedback

CERT-SCHOOL-ADMIN | 1 MONTH AGO |

[PERMALINK](#) | [REPLY](#)

Question for S3 Bucket Cross Region replication

we can Disable or delete Cross-Region Replication as per the document

<http://docs.aws.amazon.com/AmazonS3/latest/UG/cross-region-replication.html>

and there this feature is not automatic we need to enable that

